

Monthly revenue has more than doubled in a year — and ~84% of it comes from people buying for the first time

Revenue, Aug 2026 \$115K +111% vs Aug 2025	Revenue from new customers, last 12 mo 84% returning: 12–21% a month	Median time to 2nd purchase 412 days 36% ever buy again; 4.4% within 90 days	First order as share of 24-mo value 89% \$85 of \$95
---	---	---	---

- 1


Growth is real and it is volume. Orders are up 129% year over year; average order value is not (\$87 → \$80). The business is winning more customers, not bigger baskets.
- 2

Repeat purchase is slow, so growth depends on acquisition. The typical returning customer takes over a year to come back, and 70% of first-time buyers arrive through one channel, Search. That is a concentration risk, not a retention crisis.
- 3

The free-shipping test cannot be read from this data. At ~500 completed orders a quarter, only a shift of ≥8 points in the share of orders over \$100 would be detectable. The design is right; the sample is not.

Revenue grew 111% year over year on order volume — average order value did not grow

Trailing-12-month revenue is \$946K, up 70% on the prior 12 months. The typical month grows 4.5% on the one before.



+129%

completed orders, Aug 2026 vs a year earlier

\$75–\$94


range of monthly average order value over two years

4.5%

median month-over-month revenue growth, last 12 months

Roughly 8 in 10 revenue dollars every month come from first-time buyers

Returning customers are not the problem — they spend more per head (\$89 vs \$80 in Aug). There are just very few of them each month.



84%
of revenue from new customers, last 12 months

1,145
new customers in Aug 2026 — a record

+12%
spend per returning customer vs per new customer


Customers who come back take a median of 412 days — so the first order is ~89% of what a customer spends in two years

36% of customers do buy again eventually. Almost none do it quickly: 4.4% within 90 days, 16% within a year.

Cohort retention — % of each first-purchase cohort buying again in month N


Average cumulative spend per customer — 10,846 customers with 24 months of history

	M1	M2	M3	M4	M5	M6	M7	M8
Jan 2025 (430)	2.8	0.5	0.9	0.2	0.7	0.7	0.2	0.2
Feb 2025 (430)	1.4	0.9	1.2	0.9	0.2	1.2	0.5	1.2
Mar 2025 (446)	3.1	0.5	0.9	1.4	1.6	0.2	0.7	0.2
Apr 2025 (504)	2.2	1.0	0.6	0.4	1.2	0.8	1.0	0.8
May 2025 (572)	3.3	1.1	0.3	1.4	1.2	0.3	0.2	1.4
Jun 2025 (469)	3.0	0.8	0.8	0.2	0.2	0.8	0.4	0.6
Jul 2025 (521)	4.6	1.7	1.1	0.6	0.8	1.1	0.8	1.1
Aug 2025 (537)	2.6	1.1	0.9	0.6	0.7	1.1	0.7	0.7
Sep 2025 (566)	2.6	0.5	0.9	1.1	0.5	0.7	0.5	1.1
Oct 2025 (613)	2.9	1.0	1.6	1.6	2.0	1.1	1.3	1.1
Nov 2025 (633)	3.6	0.6	2.5	1.1	0.6	0.5	1.7	1.3
Dec 2025 (659)	3.2	1.4	1.1	1.5	1.2	1.8	0.9	2.1
Jan 2026 (686)	3.6	1.3	1.3	2.0	0.9	1.8	1.5	
Feb 2026 (609)	4.9	1.8	0.8	1.8	0.7	2.0		



94% of active customers don't buy again within 90 days — a window 4.6× shorter than the typical repeat cycle


Revenue and churn rise together: the churn figure describes a business that grows by adding customers, not one that is losing them. It is also the same in every acquisition channel.



94.1% → 88.5%

churn if “purchase” includes shipped and processing orders, not only ‘Complete’


12-month repeat rate by channel




Free shipping over \$100 can't be evaluated on this data: the design works, the sample is too small

Eligible orders grew 26% against 5% for the rest — a +20-point difference-in-differences that is not statistically significant ($p = 0.22$).

Difference-in-differences — completed orders 90 days before vs after



Share of orders over \$100, the metric the policy is meant to move



≥8 pts

smallest shift in the share of orders over \$100 detectable with ~500 orders a quarter (80% power)

What a real test needs: a shipping_fee field, a holdout group, cart-abandonment events, and margin net of shipping.


Recommendations and KPI framework

IN PROGRESS

This slide answers the brief's "Questions to address":

- Definitions & alternatives
- The single most important trend for leadership
- One business initiative: target segment, primary and secondary metrics
- A 5–7 metric business-health dashboard

One fact table and two joins — each checked to add columns, never rows




JOIN RULES APPLIED

- Pick the grain first.** `order_items` is one row per line item, so `orders` and `customers` are counted with `DISTINCT`.
- Count rows before and after every join.** Joining `inventory_items`: 45,373 before, 45,373 after.
- Skip joins that add nothing.** `order_items` already carries the order's `user_id` and `status` (0 mismatches), so `orders` is never joined.
- Aggregate before joining one-to-many.** Task D rolls line items up to one row per order before attaching users.

Three things about this data change the numbers, and none of them is in the brief

Order status mix by year of order — a 2019 order is still ‘Shipped’ today



1,154 rows dated in the future. The generator writes ahead of today. Every query caps its window at the earlier of the last row and today.


25% of line items are ‘Complete’, the same share every year. Status is assigned at random, so the brief’s definition samples a quarter of real orders.

5.6 pts of churn come from that sampling alone (94.1% vs 88.5%). A repeat order labelled ‘Shipped’ counts as a lost customer.


The session funnel can't measure conversion: anonymous traffic is a constant, so “conversion” climbs from 2% to 40% on its own

Every identified session ends in a purchase; no anonymous session ever does. Underneath, theLook places each order at a random date between signup and today — which is also why every calendar trend here rises into the present.

events, by year — fixed paths, fixed anonymous volume



Order timing test — 42,287 single-order customers



The funnel this data can support — user level, accounts ≥ 1 year old

- 100 signed up
- 80 placed an order
- 27 activated — first completed order
- 20 placed a second order

±0.6 pts
spread across the five acquisition channels. No channel converts or repeats differently.

AI wrote the first draft of every query; a check that could fail decided whether to keep it

WHERE IT SPED THINGS UP • Profiling the seven tables and every join key in minutes • First drafts of CTE-structured SQL for all four tasks • Generating this deck reproducibly from the query outputs	WHERE CHECKS OVERRULED IT • A Task D column was 100% or 0% by construction — caught, and the query rebuilt at monthly grain • A 100% churn month — identified as the edge of the data, not an event • “Retention is near zero” — revised once repeat timing was measured • A session funnel showing conversion rising 23× — traced to a constant denominator
WHAT I WOULD AUTOMATE NEXT • Run the QA queries on a schedule and alert when a check fails, before a dashboard refreshes • Have an LLM draft the weekly KPI narrative with every figure pulled from SQL, never computed by the model • Point text-to-SQL at curated, documented views rather than raw tables, so definitions are fixed upstream	CHECKS THAT SHIP WITH THE REPO • Row counts before and after every join • Identity assertion: new + returning = active, every month • Revenue reconciled to the status breakdown • Definition sensitivity: churn under two definitions

EXAMPLE PROMPT

“The last three months return churn near 100%. Before assuming that is real: what in the data boundary could produce it, and write the check that would tell an artifact from a genuine spike.”